ECE408/CS483/CSE408 Fall 2022

Applied Parallel Programming

Lecture 18
Atomic Operations and Histogramming

1

Course Reminders

- MP5.2 is due this week
- Project Milestone 2: Baseline Convolution Kernel
  - Due November 4th

Objective

- To understand atomic operations
    - Read-modify-write in parallel computation
    - A primitive form of "critical regions" in parallel programs
    - Use of atomic operations in CUDA
    - Why atomic operations reduce memory system throughput
    - How to avoid atomic operations in some parallel algorithms
- To learn practical histogram programming techniques
    - Basic histogram algorithm using atomic operations
    - Atomic operation throughput
    - Privatization

A Common Collaboration Pattern

- Multiple bank tellers count the total amount of cash in the safe

- Each grab a pile and count

- Have a central display of the running total

- Whenever someone finishes counting a pile, add the subtotal of the pile to the running total

- A bad outcome
    - Some of the piles were not accounted for

ECE408/CS483/ University of Illinois at Urbana-Champaign

A Common Arbitration Pattern

- Multiple customers booking air tickets
- Each
  - Brings up a flight seat map
  - Decides on a seat
  - Update the the seat map, mark the seat as taken

- A bad outcome
  - Multiple passengers ended up booking the same seat

Read-Modify-Write Operations

thread1:   Old ← Mem[x]   thread2:   Old ← Mem[x]
           New ← Old + 1   New ← Old + 1
           Mem[x] ← New   Mem[x] ← New

If Mem[x] was **initially 0**, what would the value of Mem[x] be after threads 1 and 2 have completed?

– What does each thread get in their Old variable?

The answer may vary due to data races. To avoid data races, you should use atomic operations

Timing Scenario #1

Time	Thread 1	Thread 2
1	(0) Old ← Mem[x]	
2	(1) New ← Old + 1	
3	(1) Mem[x] ← New	
4		(1) Old ← Mem[x]
5		(2) New ← Old + 1
6		(2) Mem[x] ← New

- Thread 1 Old = 0
- Thread 2 Old = 1
- Mem[x] = 2 after the sequence

Timing Scenario #2

Time	Thread 1	Thread 2
1		(0) Old ← Mem[x]
2		(1) New ← Old + 1
3		(1) Mem[x] ← New
4	(1) Old ← Mem[x]	
5	(2) New ← Old + 1	
6	(2) Mem[x] ← New	

- Thread 1 Old = 1
- Thread 2 Old = 0
- Mem[x] = 2 after the sequence

Timing Scenario #3

Time	Thread 1	Thread 2
1	(0) Old ← Mem[x]	
2	(1) New ← Old + 1	
3		(0) Old ← Mem[x]
4	(1) Mem[x] ← New	
5		(1) New ← Old + 1
6		(1) Mem[x] ← New

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Timing Scenario #4

Time	Thread 1	Thread 2
1		(0) Old ← Mem[x]
2		(1) New ← Old + 1
3	(0) Old ← Mem[x]	
4		(1) Mem[x] ← New
5	(1) New ← Old + 1	
6	(1) Mem[x] ← New	

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Atomic Operations Prevent Interleaving

thread1: Old ← Mem[x]
          New ← Old + 1
          Mem[x] ← New

                              thread2: Old ← Mem[x]
                                        New ← Old + 1
                                        Mem[x] ← New

Or

                              thread2: Old ← Mem[x]
                                        New ← Old + 1
                                        Mem[x] ← New

thread1: Old ← Mem[x]
          New ← Old + 1
          Mem[x] ← New

Without Atomic Operations

Mem[x] initialized to 0

thread1: Old ← Mem[x]

thread2: Old ← Mem[x]

New ← Old + 1

New ← Old + 1

Mem[x] ← New

Mem[x] ← New

- Both threads receive 0
- Mem[x] becomes 1

Needed When Threads Write to Same Location

When **two threads**

- **may write to** the **same memory** location,

- the program may **need atomic operations**.

Sharing is **not always easy to recognize**…

- Do two insertions into a hash table share data?

- What about two graph node updates based on all of the nodes' neighbors?

- What if nodes are on same side of bipartite graph?

What Exactly is "Atomic?"

To a high-energy photon, atoms are not.

Atomicity is ALWAYS with respect to something.

Two **sections of code**

- **that execute atomically with respect to one another**

- **appear** to the software **as though**

- the programs' **execution did not interleave** at all.

What Can Go Wrong?

Common failure mode:

- **Programmer *thinks* operations** are **independent**.
- Hasn't considered input data for which they are not.
- Or another programmer reuses code without understanding assumptions that imply independence.

Also: **atomicity does not constrain relative order**.

Implementing Atomic Operations

- Many ISAs offer synchronization primitives,
  - **instructions with** one (or more) **address operands**
  - **that execute atomically with respect to one another** when used on the same address.
- Mostly read, modify, write operations
  - Bit test and set
  - Compare and swap / exchange
  - Swap / exchange
  - Fetch and increment / add

Atomicity Enforced by Microarchitecture

When synchronization primitives execute,

- **hardware ensures** that **no other** thread
- **accesses** the location **until** the operation is **complete**.

Other threads that access the location

- are typically stalled or held in a queue until their turn.
- **Threads perform atomic operations serially**.

Atomic Operations in CUDA

- Function calls that are translated into single ISA instructions (a.k.a. *intrinsics*)

    - Atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)

    - Read CUDA C programming Guide for more details

- Atomic Add

    int atomicAdd(int **address**, int **val**);*

    reads the 32-bit word **old** pointed to by **address** in global or shared memory, computes **(old + val)**, and stores the result back to memory at the same address. The function returns **old**.

More Atomic Adds in CUDA

- Unsigned 32-bit integer atomic add

  unsigned int atomicAdd(unsigned int address, unsigned int val);*

- Unsigned 64-bit integer atomic add

  unsigned long long int atomicAdd(unsigned long long int address, unsigned long long int val);*

- Single-precision floating-point atomic add (capability > 2.0)

  – float atomicAdd(float* address, float val);

Building synchronization with atomics

- How would we build `__syncthreads()` for block?
- How would we create `__syncthreads()` for entire grid?
  - And why would this not be a good idea?

- How would we create a critical section? I.e., one thread per block executing a particular section of code?

- How would we create a critical section per grid?
  - Why doesn't this have the same issue as `__syncthreads()` for grid?

Atomic Compare and Swap (CAS)

```
Bool atomicCAS(int *address, int old, int new)
{
    if (*address != old)
        return false;


    *address = new;


    return true;
}
```

Atomic Add using CAS

```
int atomicAdd(int* address, int value)
{
    Bool done = false;

    while (!done) {
        old_v = *address;
        done = atomicCAS(address, old_v, old_v + value);
    }

    return old_v;
}
```

Histogramming

- A method for extracting notable features and patterns from large data sets
  - Feature extraction for object recognition in images
  - Fraud detection in credit card transactions
  - Correlating heavenly object movements in astrophysics
  - …

- Basic histograms - for each element in the data set, use the value to identify a "bin" to increment

A Histogram Example

- In sentence "Programming Massively Parallel Processors" build a histogram of frequencies of each letter

- A(4), C(1), E(1), G(1), …

- How do you do this in parallel?

Iteration #1 – 1ˢᵗ letter in each section
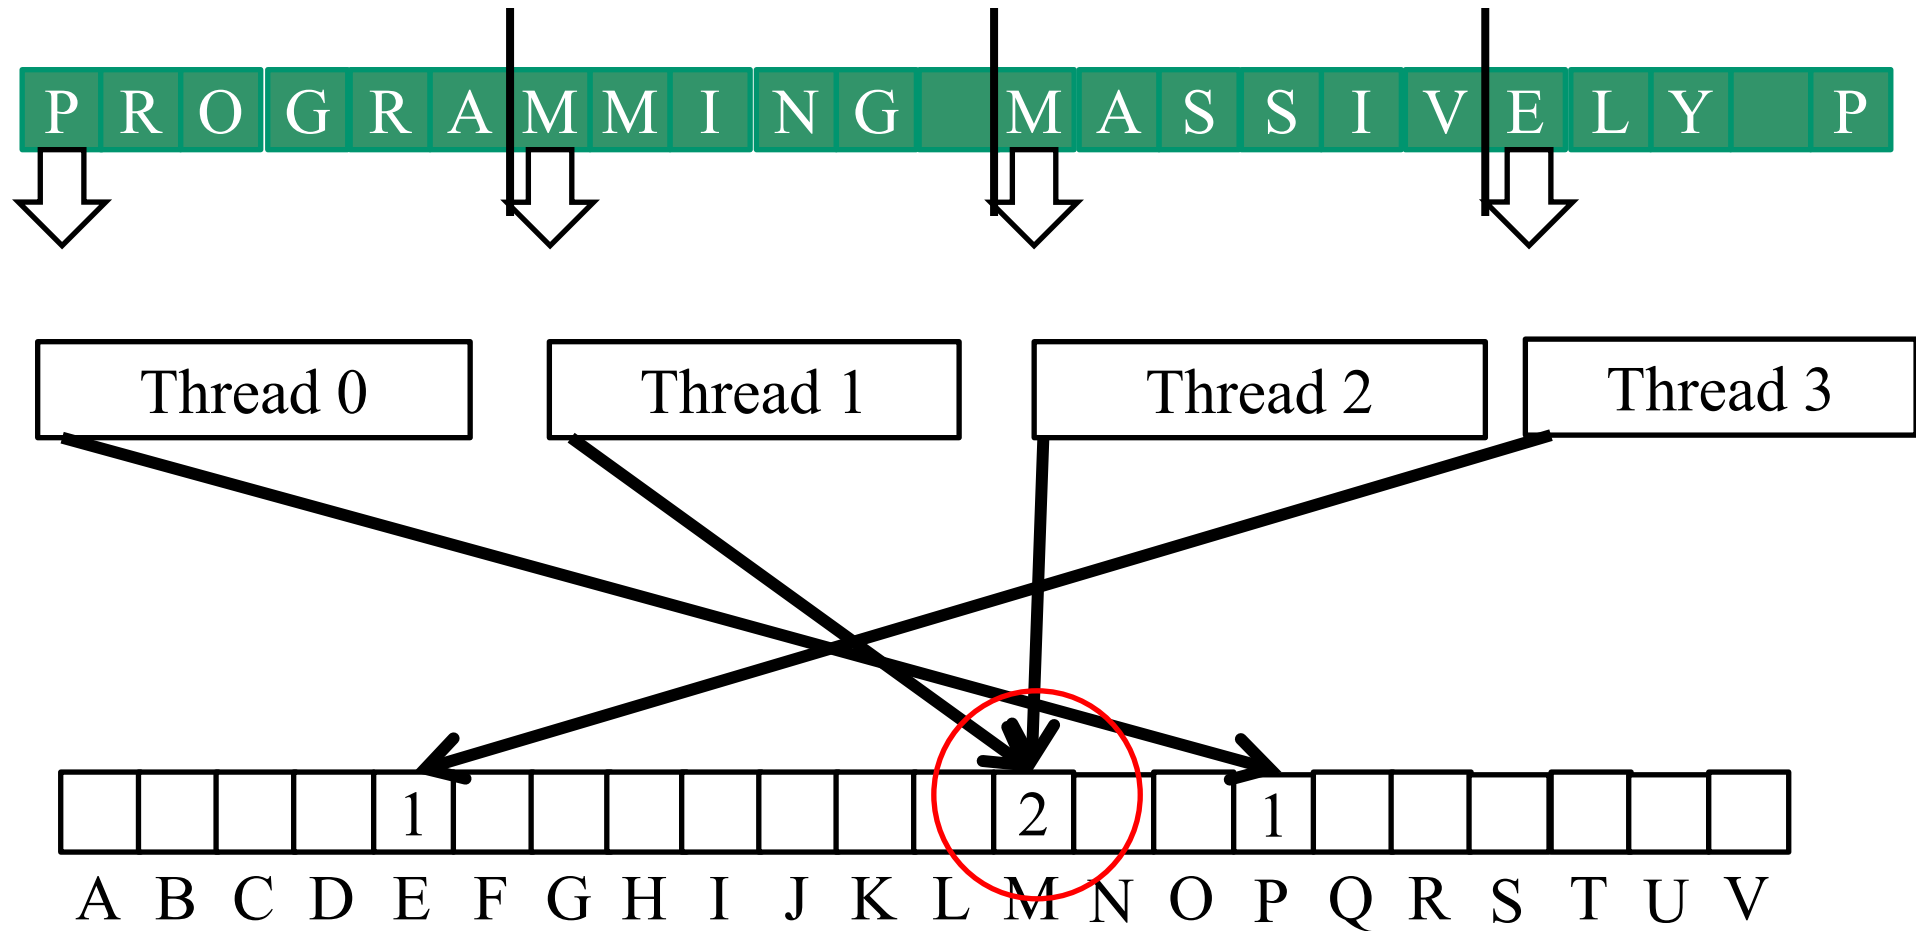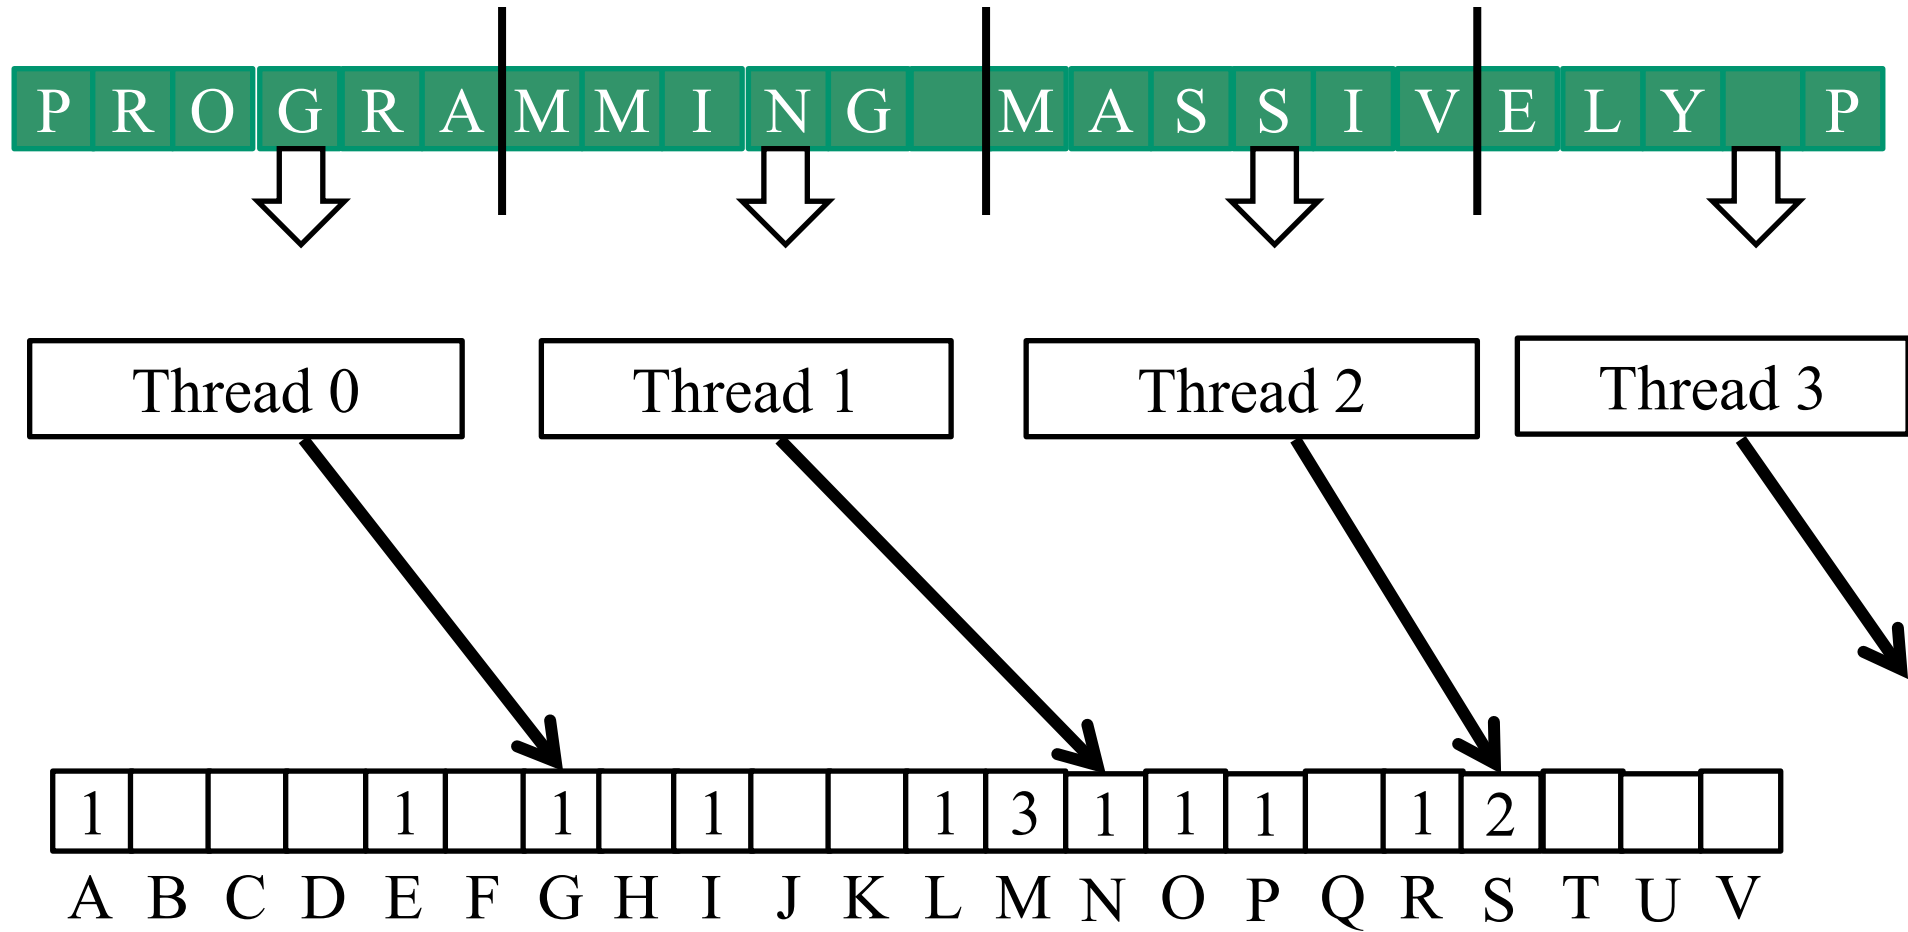

Atomic operation enforces correct update

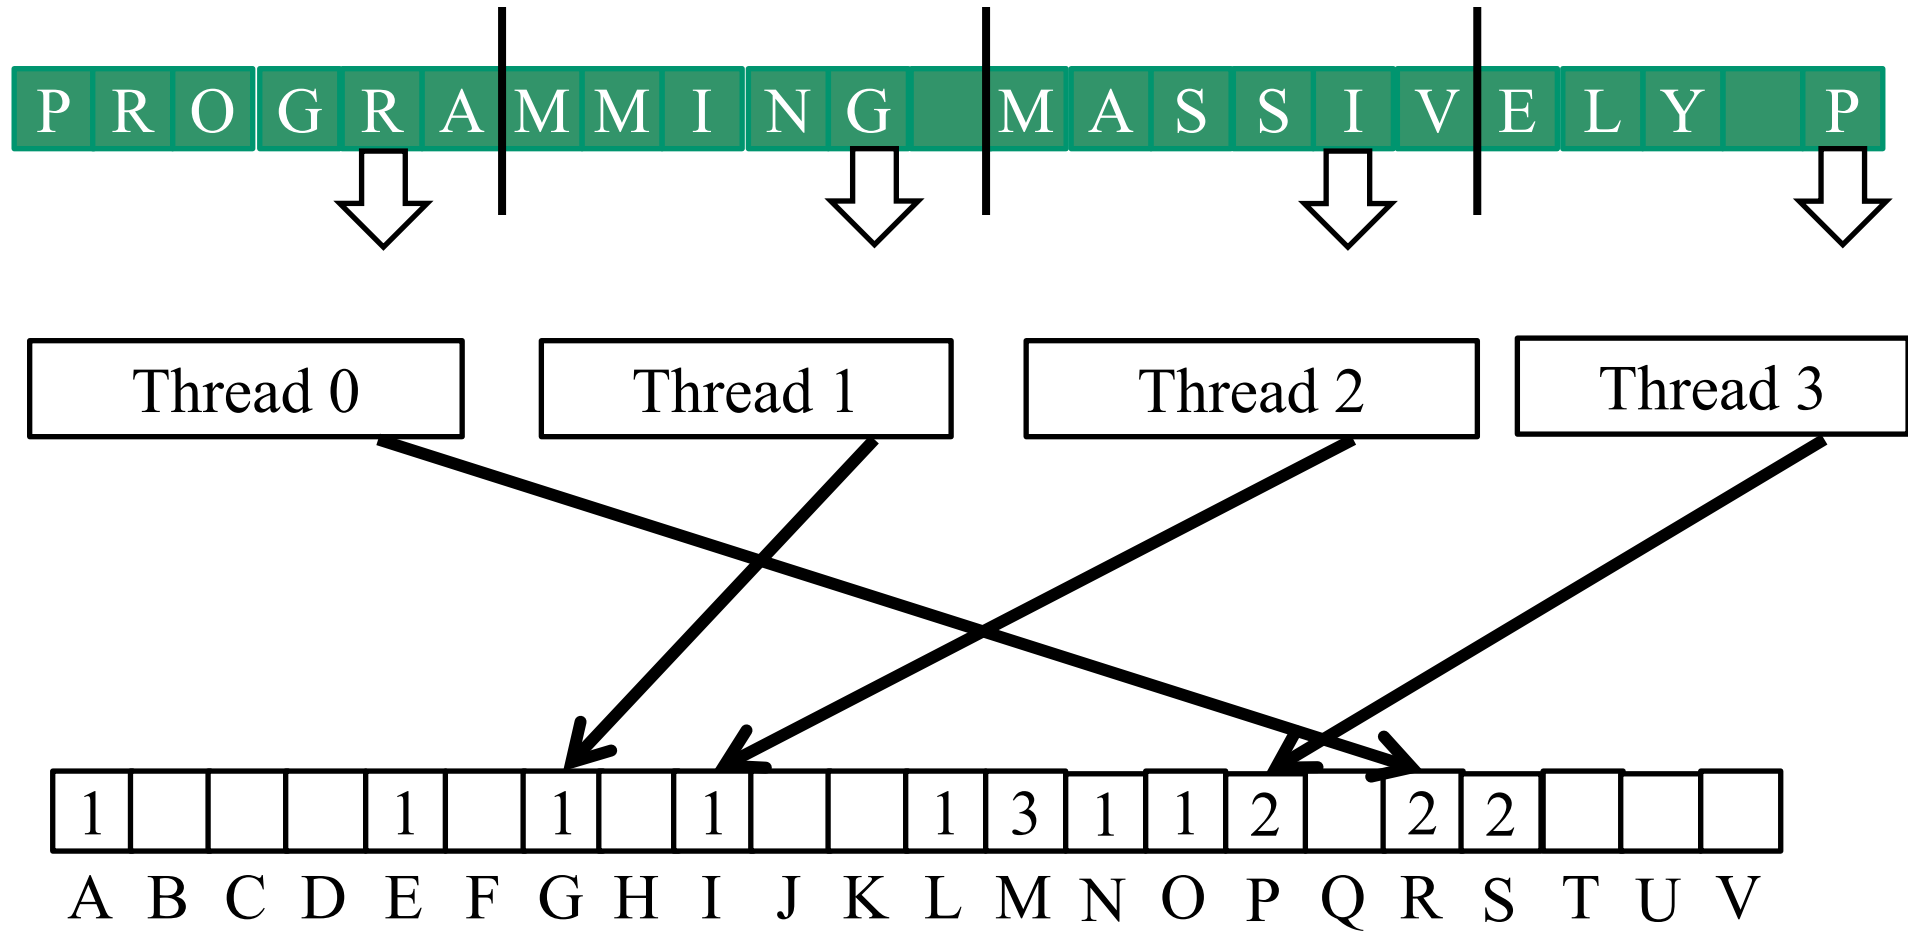# Iteration #2 – 2nd letter in each section
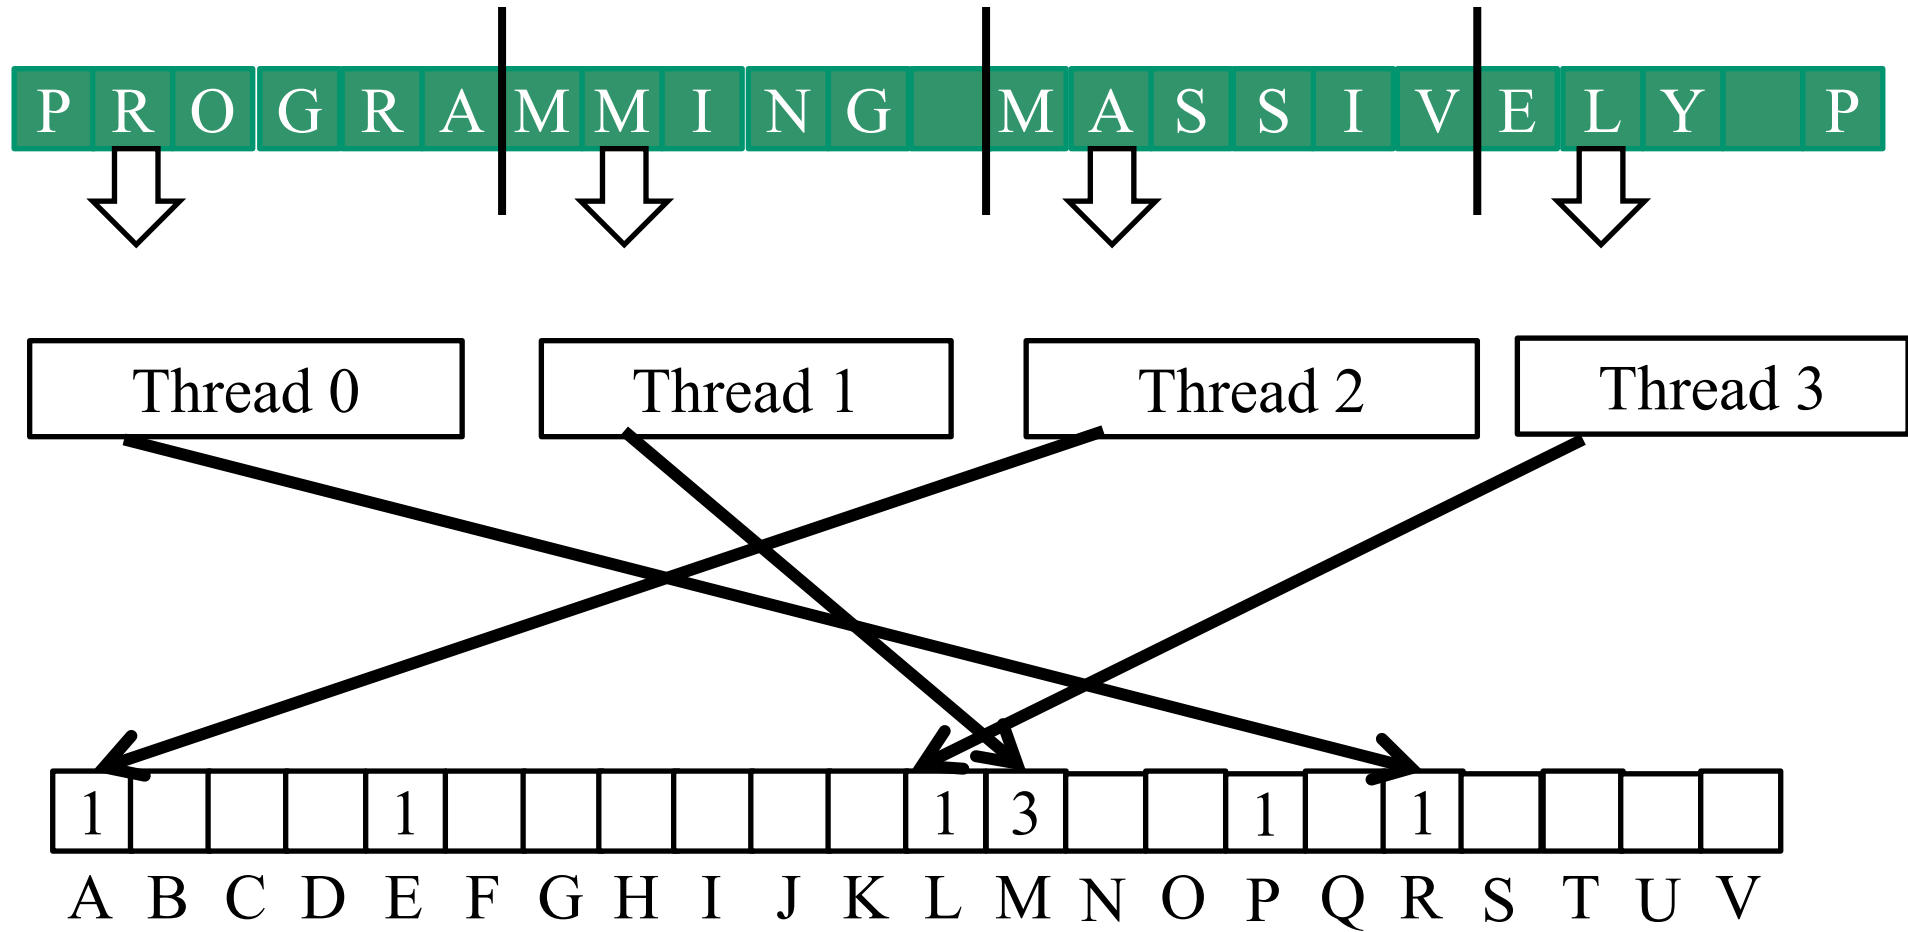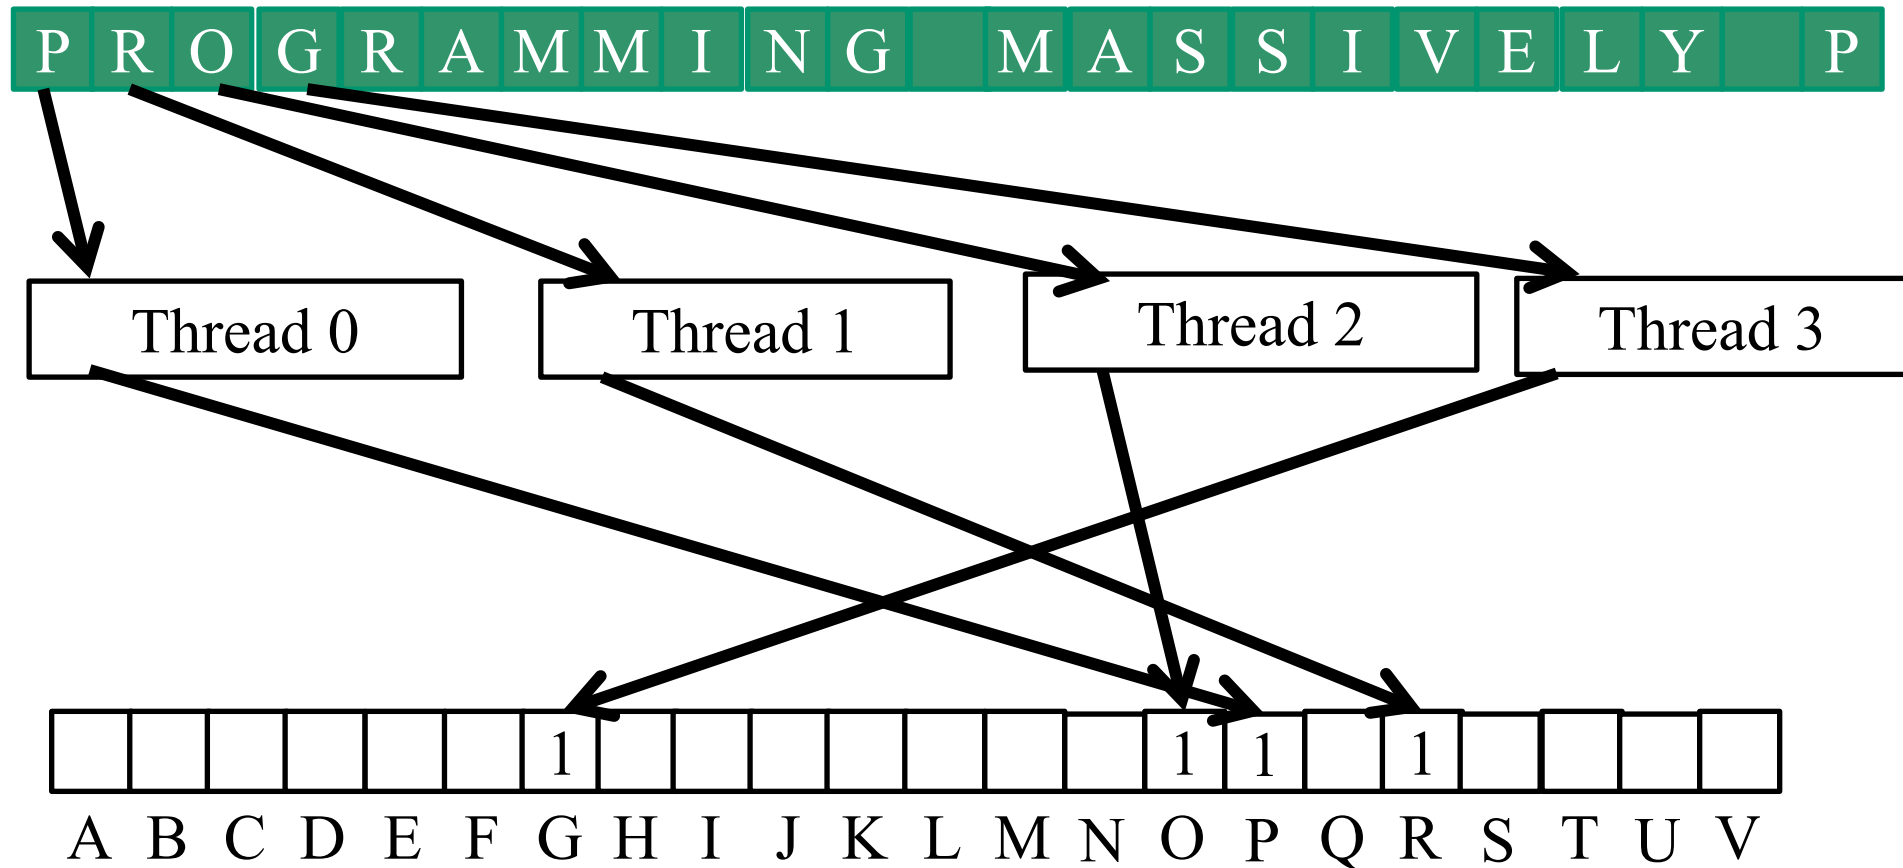
Iteration #3
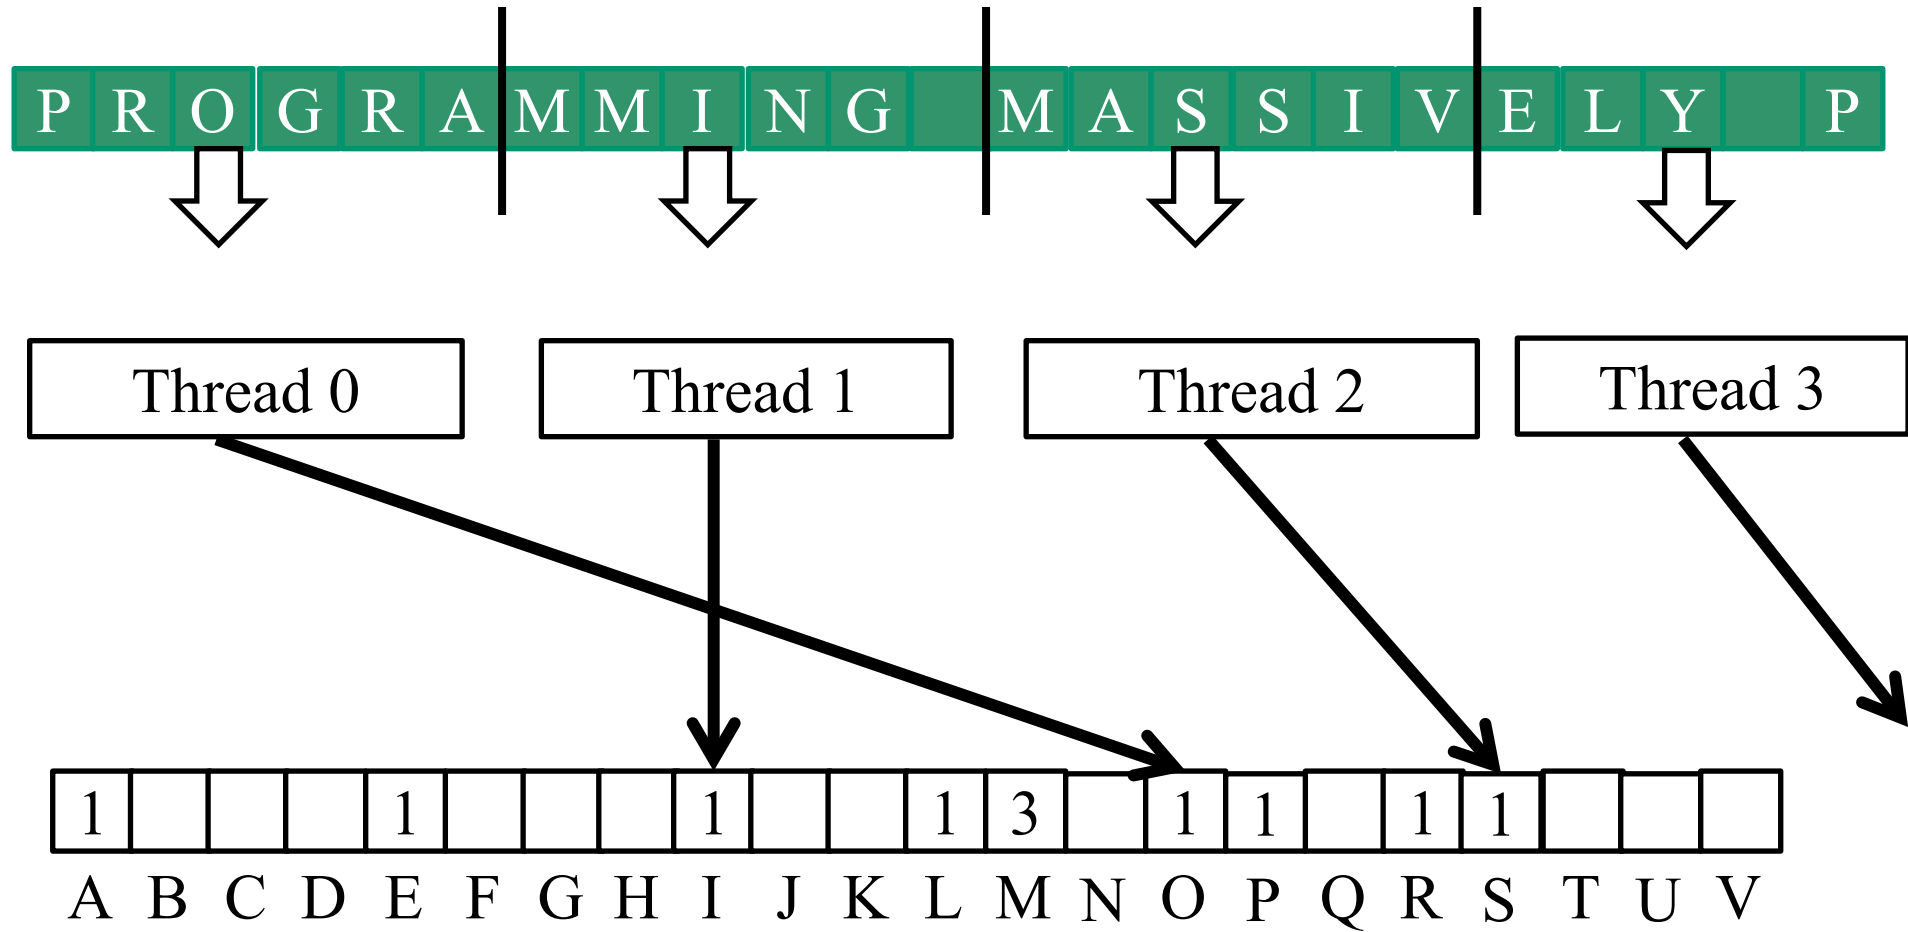
Iteration #4

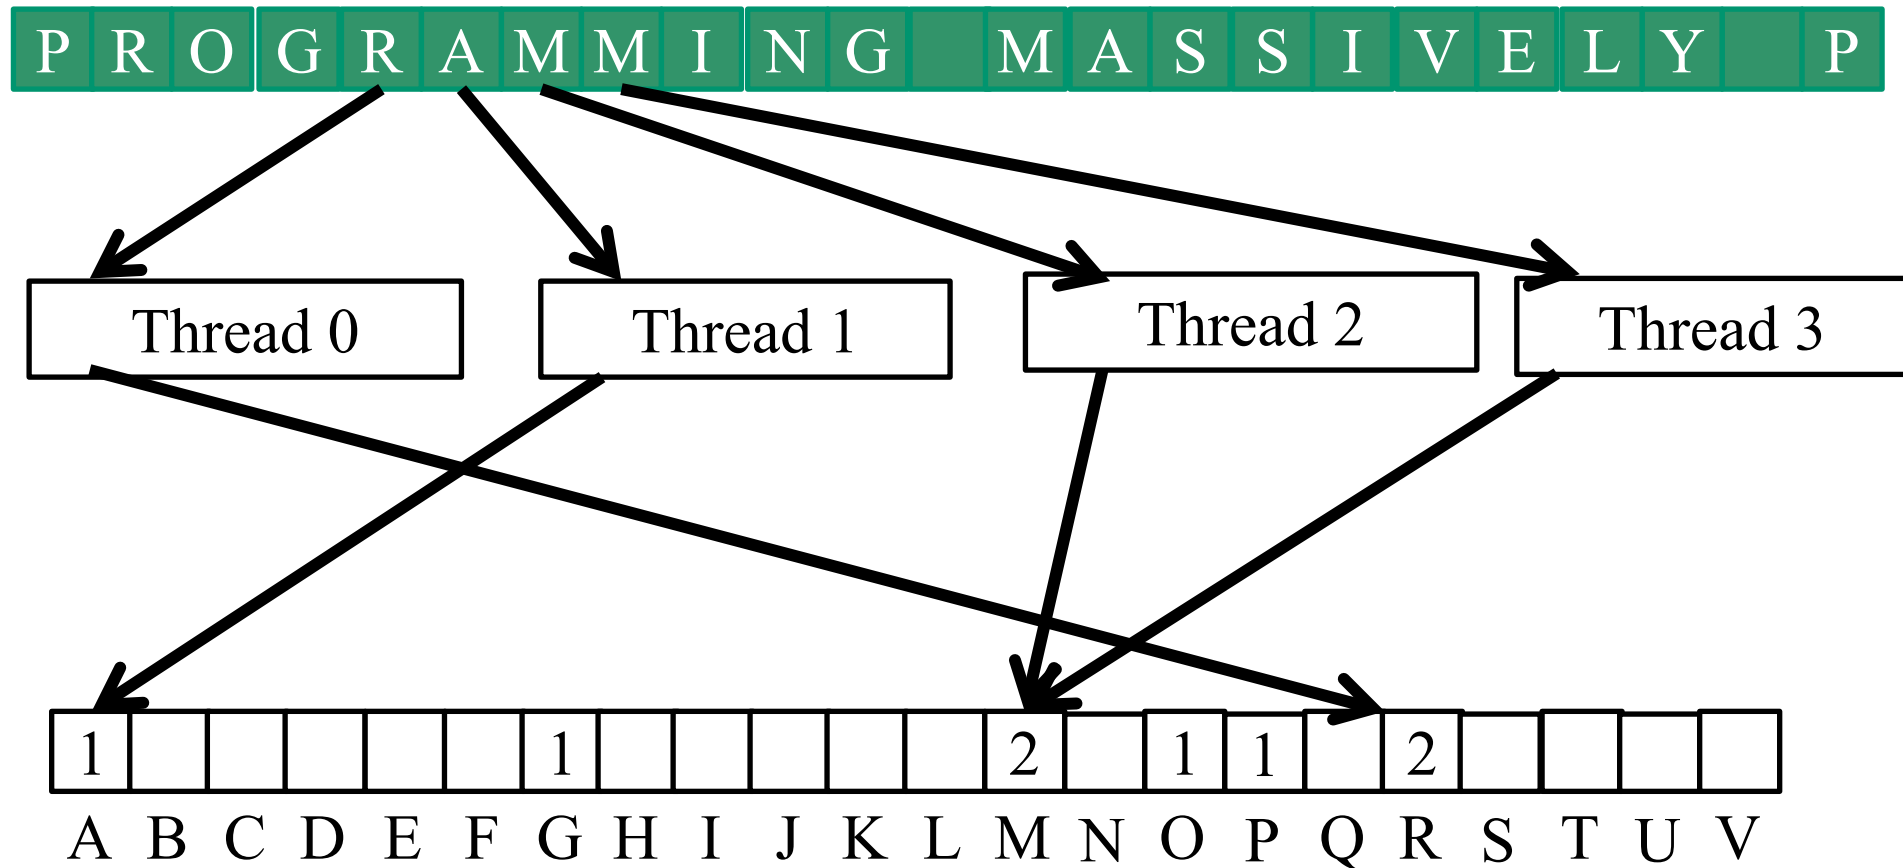# Iteration #5

A better approach

- Reads from the input array are not coalesced
  - Assign inputs to each thread in a strided pattern
  - Adjacent threads process adjacent input letters

| P | R | O | G | R | A | M | M | I | N | G | | M | A | S | S | I | V | E | L | Y | | P |

Thread 0   Thread 1   Thread 2   Thread 3

| | | | | | | 1 | | | | | | | | 1 | 1 | | 1 | | | | |

A B C D E F G H I J K L M N O P Q R S T U V

Iteration 2

- All threads move to the next section of input

A Histogram Kernel

- The kernel receives a pointer to the input buffer
- Each thread process the input in a strided pattern

```
__global__
void histo_kernel(unsigned char *buffer,
                    long size, unsigned int *histo)
{
    int i = threadIdx.x + blockIdx.x * blockDim.x;


    // stride is total number of threads
    int stride = blockDim.x * gridDim.x;
```

More on the Histogram Kernel

```
// All threads in the grid collectively handle
// blockDim.x * gridDim.x consecutive elements

while (i < size) {
    atomicAdd( &(histo[buffer[i]]), 1);
    i += stride;
}
}
```
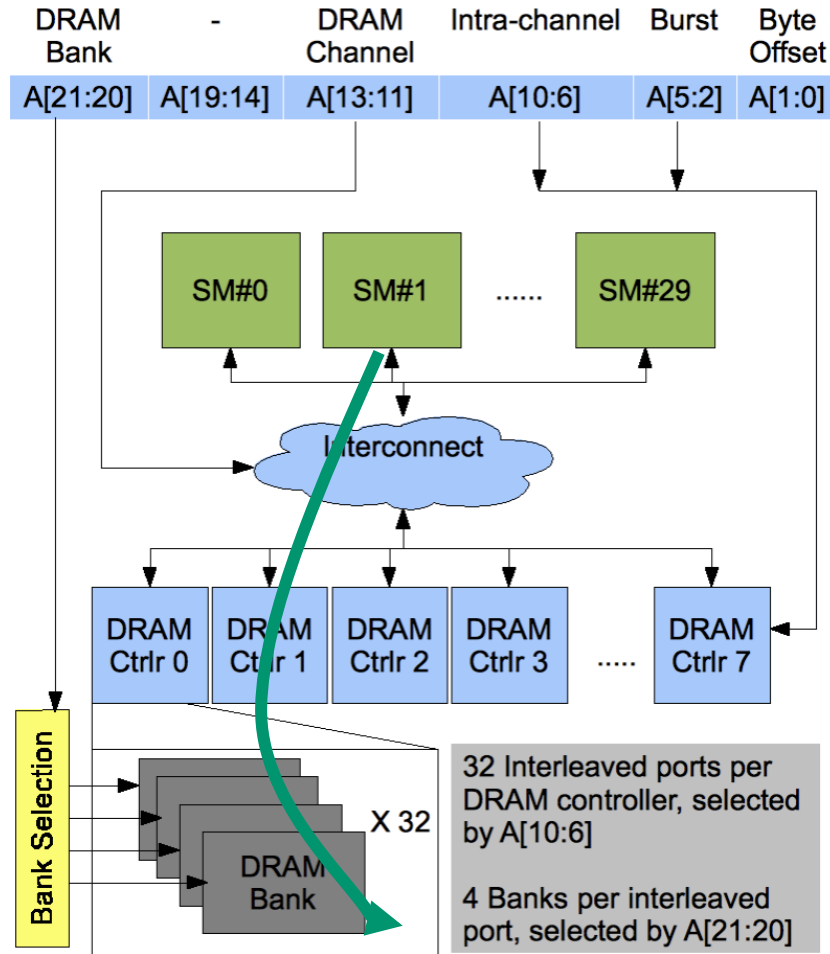
Atomic Operations on DRAM
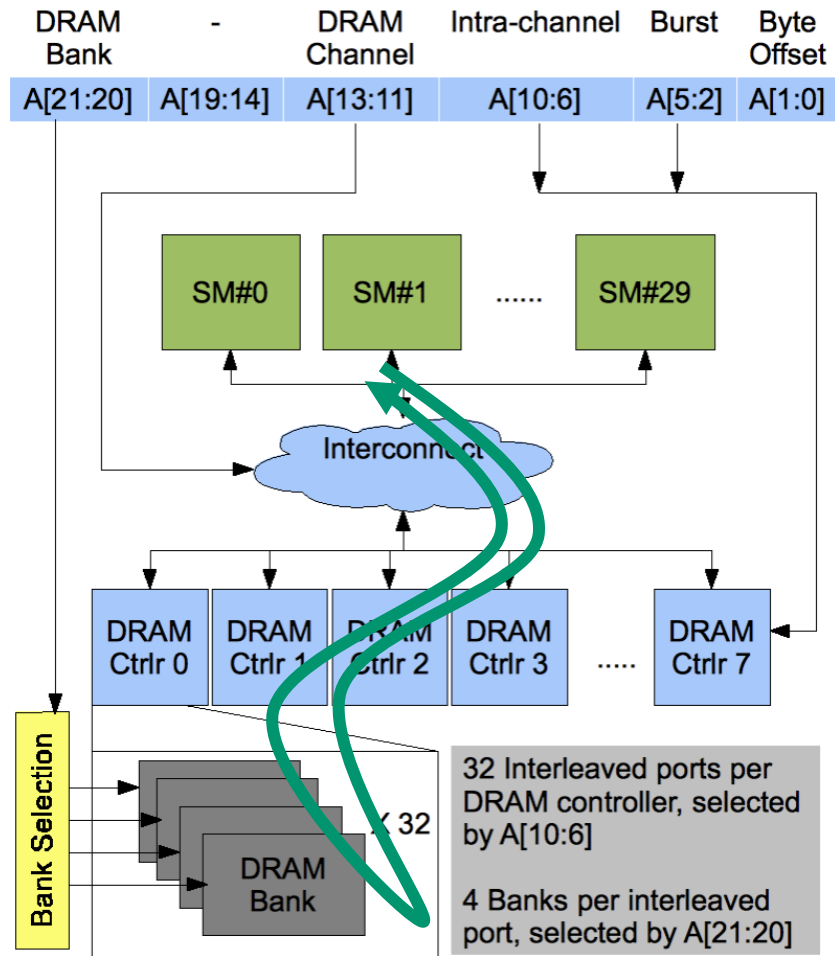
- An atomic operation starts with a read, with a latency of a few hundred cycles
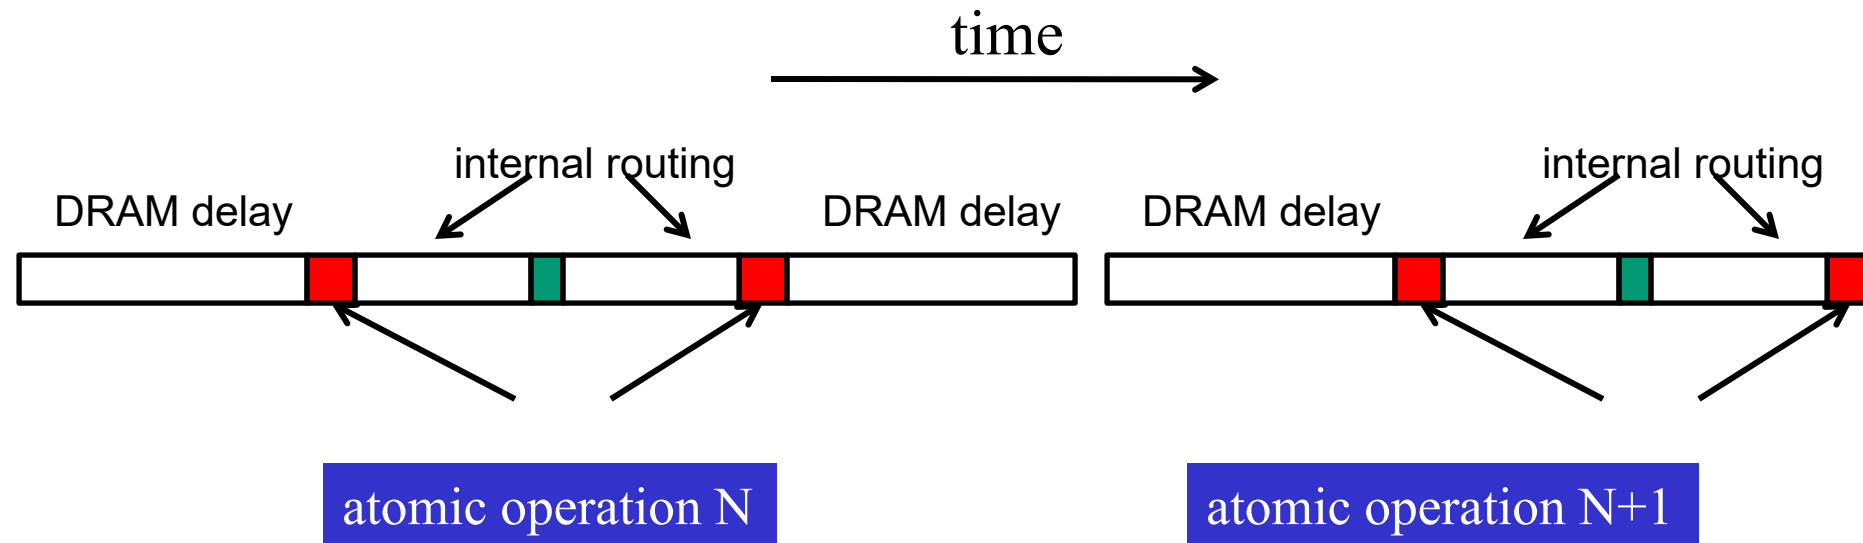
Atomic Operations on DRAM

- An atomic operation starts with a read, with a latency of a few hundred cycles
- The atomic operation ends with a write, with a latency of a few hundred cycles
- During this whole time, no one else can access the location

Atomic Operations on DRAM

- Each Load-Modify-Store has two full memory access delays
  - All atomic operations on the same variable (RAM location) are serialized

time

internal routing   internal routing

DRAM delay   DRAM delay   DRAM delay

atomic operation N   atomic operation N+1

# Latency determines throughput of atomic operations

- Throughput of an atomic operation is the rate at which the application can execute an atomic operation on a particular location.

- The rate is limited by the total latency of the read-modify-write sequence, typically more than 1000 cycles for global memory (DRAM) locations.

- This means that if many threads attempt to do atomic operation on the same location (contention), the memory bandwidth is reduced to < 1/1000!

You may have a similar experience in supermarket checkout

- Some customers realize that they missed an item after they started to check out

- They run to the isle and get the item while the line waits
  - The rate of check is reduced due to the long latency of running to the isle and back.

- Imagine a store where every customer starts the check out before they even fetch any of the items
  - The rate of the checkout will be 1 / (entire shopping time of each customer)

Hardware Improvements

- Atomic operations on L2 cache
  - medium latency, but still serialized
  - Global to all blocks
  - "Free improvement" on Global Memory atomics

time →

internal routing

atomic operation N   **atomic operation N+1**

Hardware Improvements

- Atomic operations on Shared Memory
  - Very short latency, but still serialized
  - Private to each thread block
  - Need algorithm work by programmers (more later)

Atomics in Shared Memory Requires Privatization

- Create private copies of the histo[] array for each thread block

```
__global__
void histo_kernel(unsigned char *buffer,
                        long size, unsigned int *histo) {
  __shared__ unsigned int histo_private[256];
    // warning: this will not work correctly if there are fewer than 256 threads!
  if (threadIdx.x < 256)
      histo_private[threadIdx.x] = 0;

  __syncthreads();
```

Build Private Histogram

```
int i = threadIdx.x + blockIdx.x * blockDim.x;

// stride is total number of threads
int stride = blockDim.x * gridDim.x;

while (i < size) {
    atomicAdd( &(private_histo[buffer[i]]), 1);
    i += stride;
}
```

$\log(n)$

Build Final Histogram

```
// wait for all other threads in the block to finish
__syncthreads();

if (threadIdx.x < 256)
    atomicAdd( &(histo[threadIdx.x]),
                    private_histo[threadIdx.x] );

}
```

More on Privatization

- Privatization is a powerful and frequently used techniques for parallelizing applications

- The operation needs to be associative and commutative
  - Histogram add operation is associative and commutative

- The histogram size needs to be small
  - Fits into shared memory

- What if the histogram is too large to privatize?

ANY MORE QUESTIONS
READ CHAPTER 9