



ECE408/CS483/CSE408 Fall 2022

Applied Parallel Programming

## Lecture 21

# GPU as part of the PC Architecture

# Course Reminders


- MP 5.1/5.2
  - We are grading it now
- MP 6
  - due next week
- Project PM 1
  - Grading now, check your grade and email TA if not graded by Friday
- Project PM 2
  - due this Friday

	Lecture 21	Nov. 3	GPU as a part of the PC Architecture
	PM 2	Nov. 4	Baseline GPU Convolution Kernel (GitHub, report)
12		Nov. 8	No lecture due to a holiday
	Lecture 22	Nov. 10	Task parallelism and asynchronous data transfer
	Lab 6	Nov. 11	Histogramming
13	Lecture 23	Nov. 15	Other acceleration APIs: OpenACC, OpenCL, OneAPI, Hip
	Lecture 24	Nov. 17	Guest Lecture: Ben Sander (AMD) 2:00-3:20pm US CT, on-line via zoom
	Lab 7	Nov. 18	Sparse Matrix Multiply
	Fall break		
14	Lecture 25	Nov. 29	Guest Lecture: James Reinders (Intel) 9:30-10:50am US Central, on-line via zoom
	Lecture 26	Dec. 1	Generalizing Parallelism and Course Retrospective
	PM 3	Dec. 2	GPU Convolution Kernel Optimizations
15	Exam 2	Dec. 6	<b>Midterm 2</b>

# Objectives

- To understand the impact of data transfers on performance when using a GPU as a co-processor
  - speeds and feeds of traditional CPU
  - speeds and feeds when employing a GPU
- To develop a knowledge base for performance tuning for modern GPUs

# Review: Canonical CUDA Program Structure

- Global variables declaration
  - Kernel functions
    - `__global__ void kernelOne(...)`
  - Main ()      // host code
    - allocate memory space on the device – `cudaMalloc(&d_GlbIVarPtr, bytes )`
    - transfer data from host to device – `cudaMemcpy(d_GlbIVarPtr, h_Gl...)`
    - execution configuration setup
    - kernel call – `kernelOne<<<execution configuration>>>( args... );`
    - transfer results from device to host – `cudaMemcpy(h_GlbIVarPtr,...)`
    - optional: compare against golden (host computed) solution
- 
- repeat  
as needed

# Bandwidth:

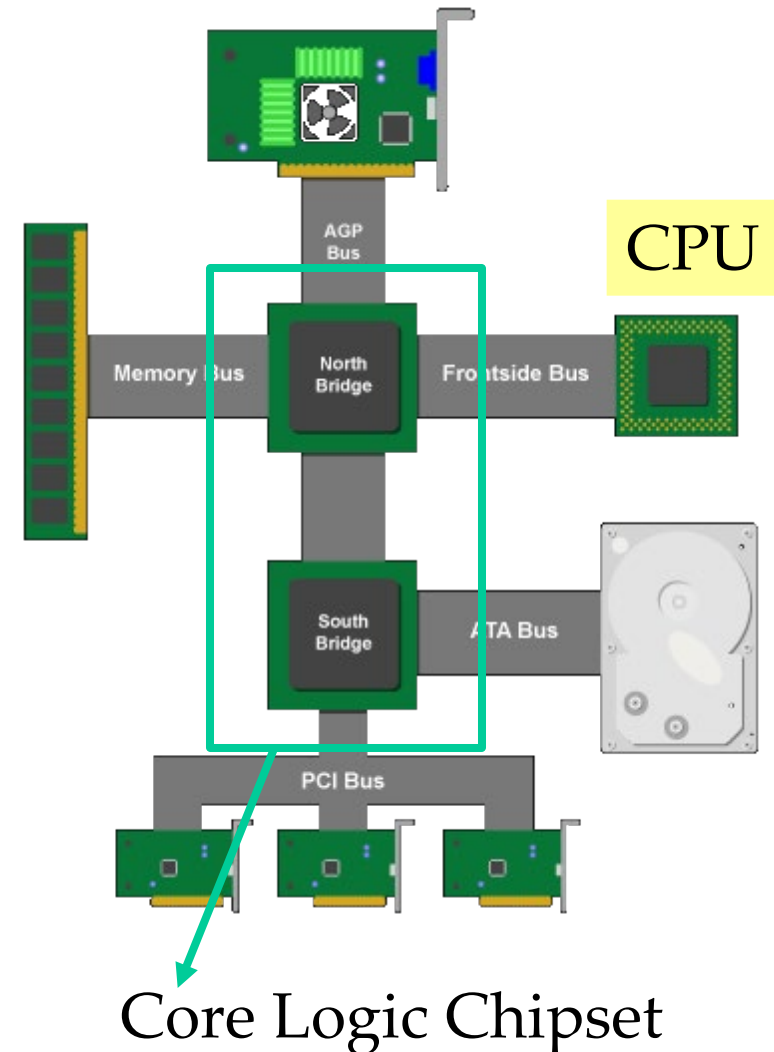
## The Gravity of Modern Computer Systems

**Bandwidth** between key components ultimately **dictates system performance**

- **Especially for GPUs** processing large amounts of data.
- Tricks like buffering, reordering, caching can temporarily defy the rules in some cases.
- Ultimately, performance falls back to what the “speeds and feeds” dictate.


# Classic (Historical) PC Architecture

- Northbridge connects 3 components that must communicate at high speed
  - CPU, DRAM, video
  - Video needs first-class access to DRAM
  - Previous NVIDIA cards are connected to AGP, up to 2 GB/s transfers
- Southbridge serves as a concentrator for slower I/O devices



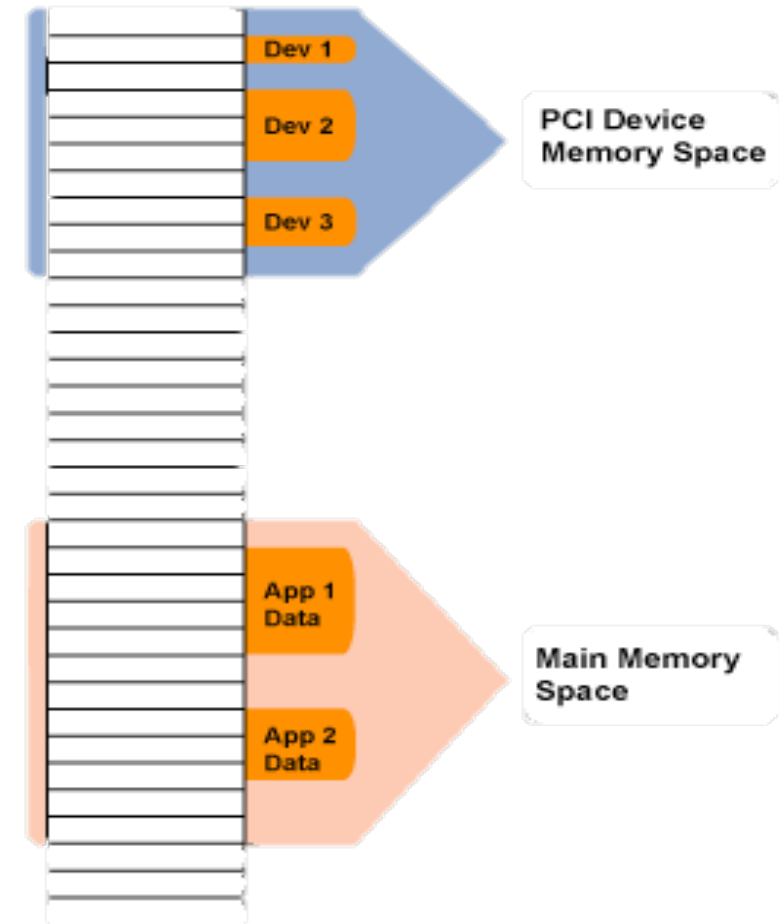
# (Original) PCI Bus Specification

- Connected to the South Bridge
  - Originally 33 MHz, 32-bit wide, 132 MB/second peak transfer rate
  - Later, 66 MHz, 64-bit, 528 MB/second peak
  - Upstream bandwidth remain slow for device ( $\sim 256\text{MB/s}$  peak)
  - **Shared bus with arbitration**
    - Winner of arbitration becomes bus master and can connect to CPU or DRAM through the southbridge and northbridge



# PCI as Memory Mapped I/O

- PCI device registers are mapped into the CPU's physical address space
  - Accessed through loads/stores (kernel mode)
- Addresses are assigned to the PCI devices at boot time
  - All devices listen for their addresses






# PCI Express (PCIe)

## **switched, point-to-point connection**

- each card has dedicated “link” to the central switch, with no arbitration
- packet switches: messages form virtual channel
- prioritized packets for QoS (such as for real-time video streaming)

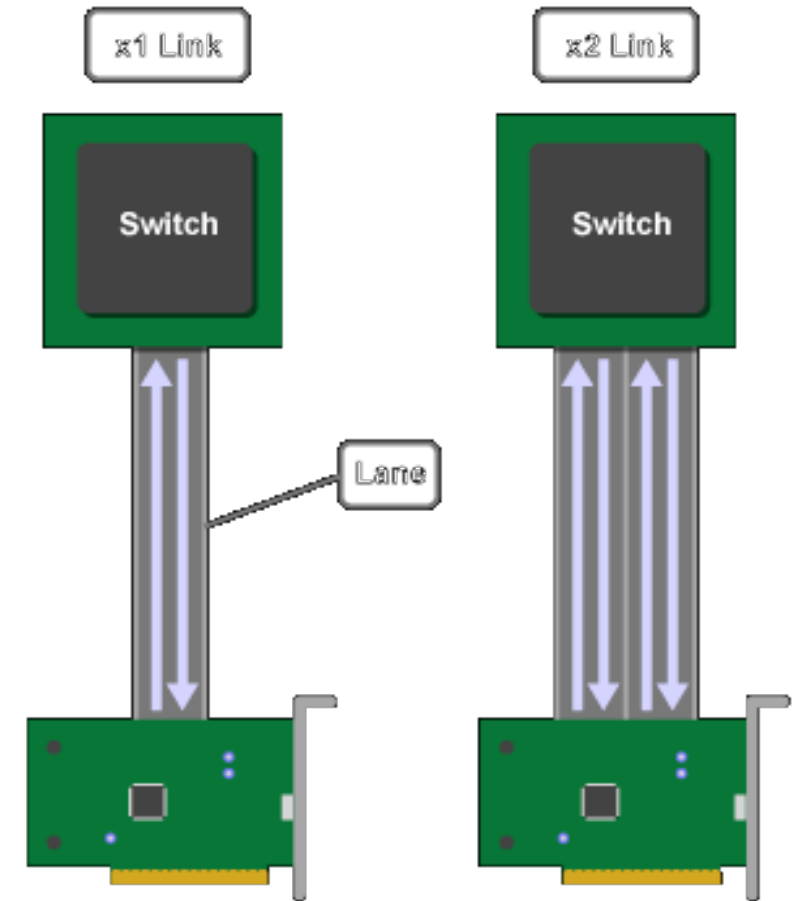


# PCIe Generations

- Within a generation, number of lanes in a link can be scaled
  - using distinct physical channels (more bits / wider transfers)
  - $\times 1, \times 2, \times 4, \times 8, \times 16, \times 32, \dots$
- Each new generation aims to double the speed
  - Current generation is PCIe 5.0, however it is supported only on a very limited set of systems, e.g., IBM Power10
    - 32GT/s
  - PCIe 4.0 is supported on modern AMD, Intel, and IBM systems
  - However, PCIe Gen. 3 is still very widely used

# PCIe Gen 3 Links and Lanes

- Each link consists of one or more lanes
  - Each lane is 1-bit wide (4 wires, each 2-wire pair can transmit 8Gb/s in one direction)
    - 2-wire pair is used for differential signaling
    - Upstream and downstream simultaneous and symmetric
  - Each Link can combine 1, 2, 4, 8, 12, 16 lanes- x1, x2, etc.
- Each byte data is **128b/130b** encoded into 130 bits with equal number of 1's and 0's; net data rate 7.8768 GB/s per lane each way.
  - Thus, the net data rates are 985 MB/s (x1) 1.97 GB/s (x2), 3.94 GB/s (x4), 7.9 GB/s (x8), 15.8 GB/s (x16), each way



# Foundation: 8/10 bit encoding

- Goal is to maintain DC balance while have sufficient state transition for clock recovery
- The difference of 1s and 0s in a 20-bit stream should be  $\leq 2$
- There should be no more than 5 consecutive 1s or 0s in any stream
- 00000000, 00000111, 11000001 bad
- 01010101, 11001100 good
- Find 256 good patterns among 1024 total patterns of 10 bits to encode an 8-bit datum
- a 20% overhead

# Current: 128/130 bit encoding

- Same goal: maintain DC balance while have sufficient state transition for clock recovery
- 1.5% overhead instead of 20%
- Scrambler function: long runs of 0s, 1s vanishingly small
- Instead of guaranteed run length of 8/10b
- At least one bit shift every 66 bits

# Patterns Contain Many 0s and 1s

**A question** for fun:

- if we need  **$2^{128}$  code words**
- **chosen from** all  $2^{130}$  **130-bit patterns**
- **how many 0s/1s** must we consider including?

**Answer: 63-67 (of either type)**


Thus 128b/130b code words are pretty well-balanced,  
and have lots of 0-1 transitions (for clock recovery).

# Recent PCIe PC Architecture

PCIe forms the interconnect backbone within PC.

Northbridge and Southbridge are PCIe switches.

Source: Jon Stokes, PCI Express: An Overview (<http://arstechnica.com/articles/paedia/hardware/pcie.ars>)




# Recent PCIe PC Architecture

How is PCI supported?

- Need a PCI-PCIe bridge, which is
- sometimes included as part of Southbridge, or
- can add as a separate PCIe I/O card.

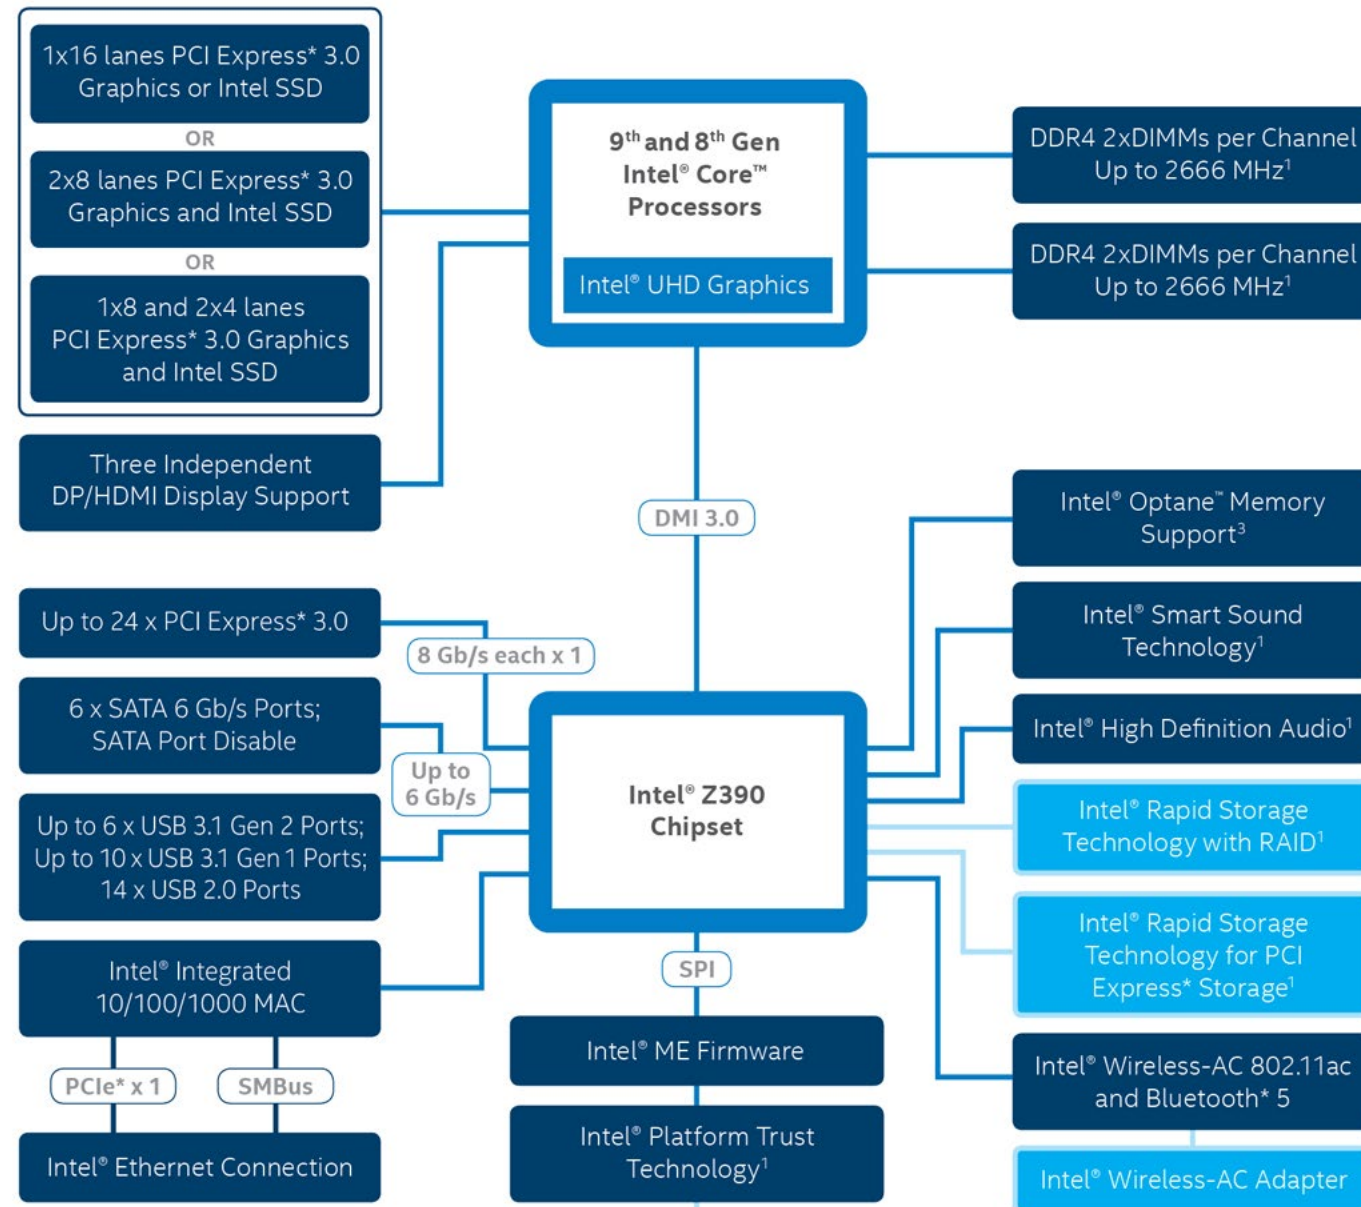
Current systems integrate PCIe controllers directly on chip with CPU.

Source: Jon Stokes, PCI Express: An Overview (<http://arstechnica.com/articles/paedia/hardware/pcie.ars>)

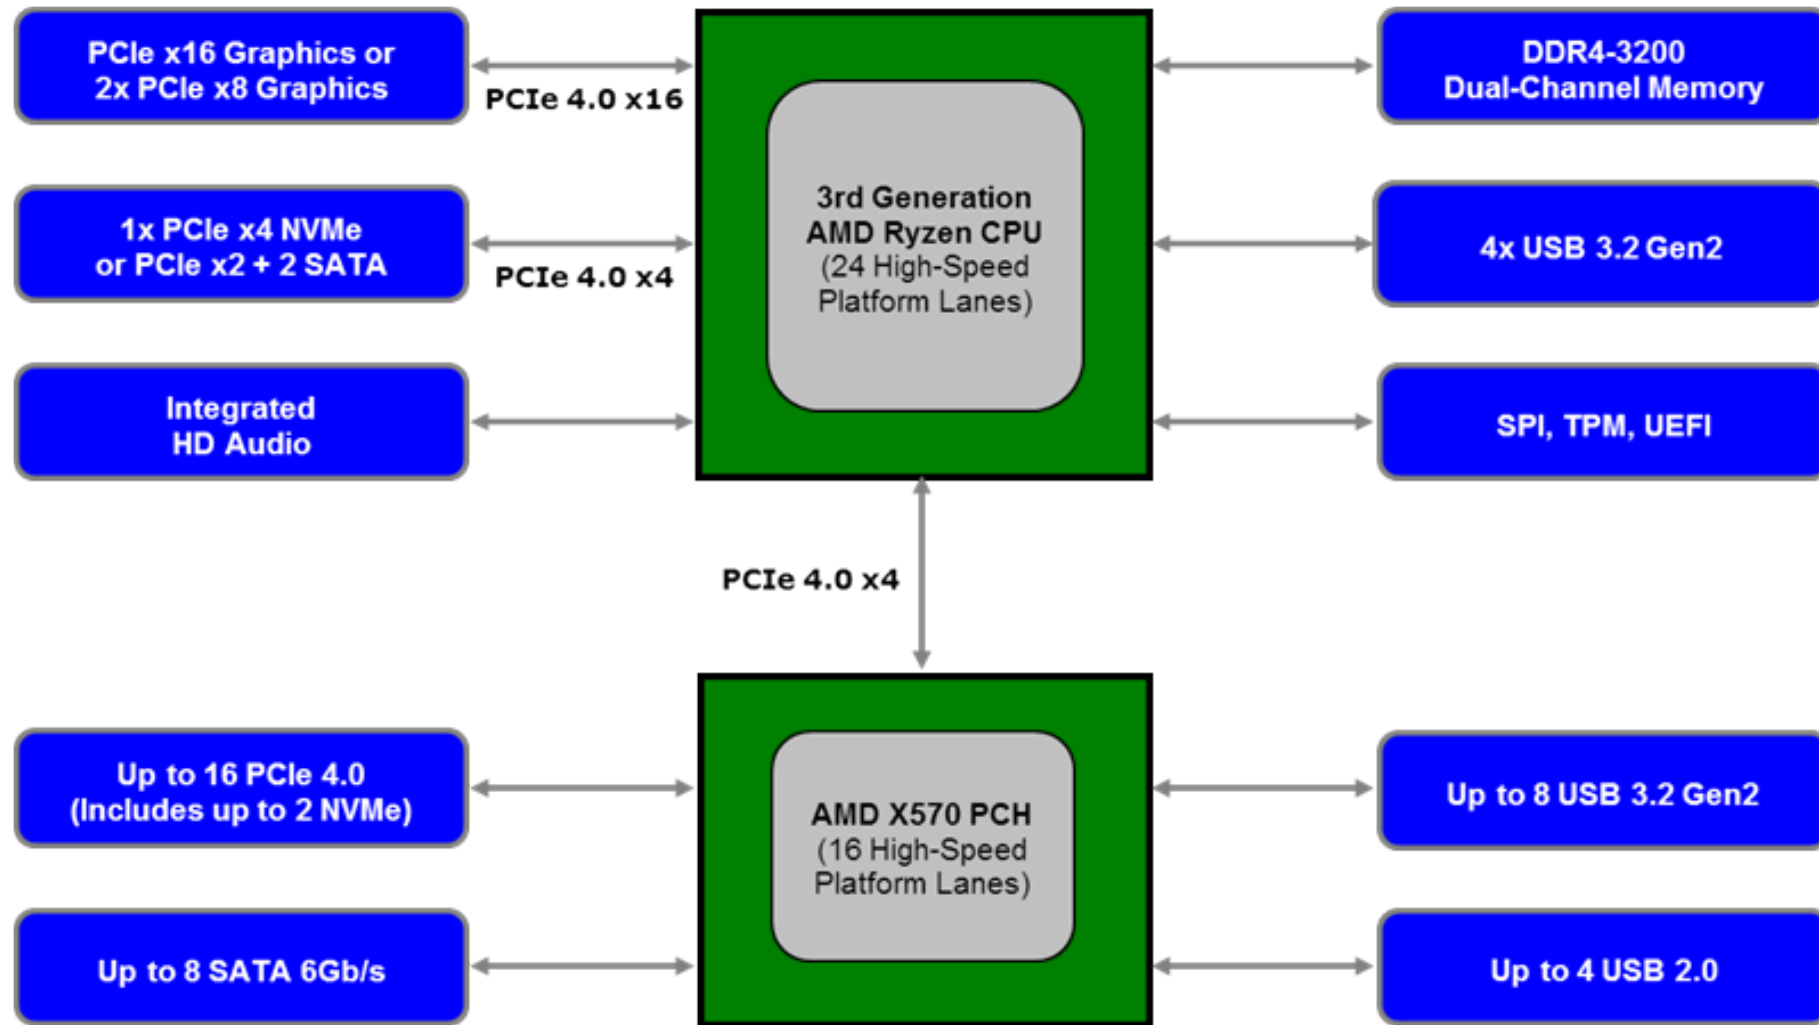




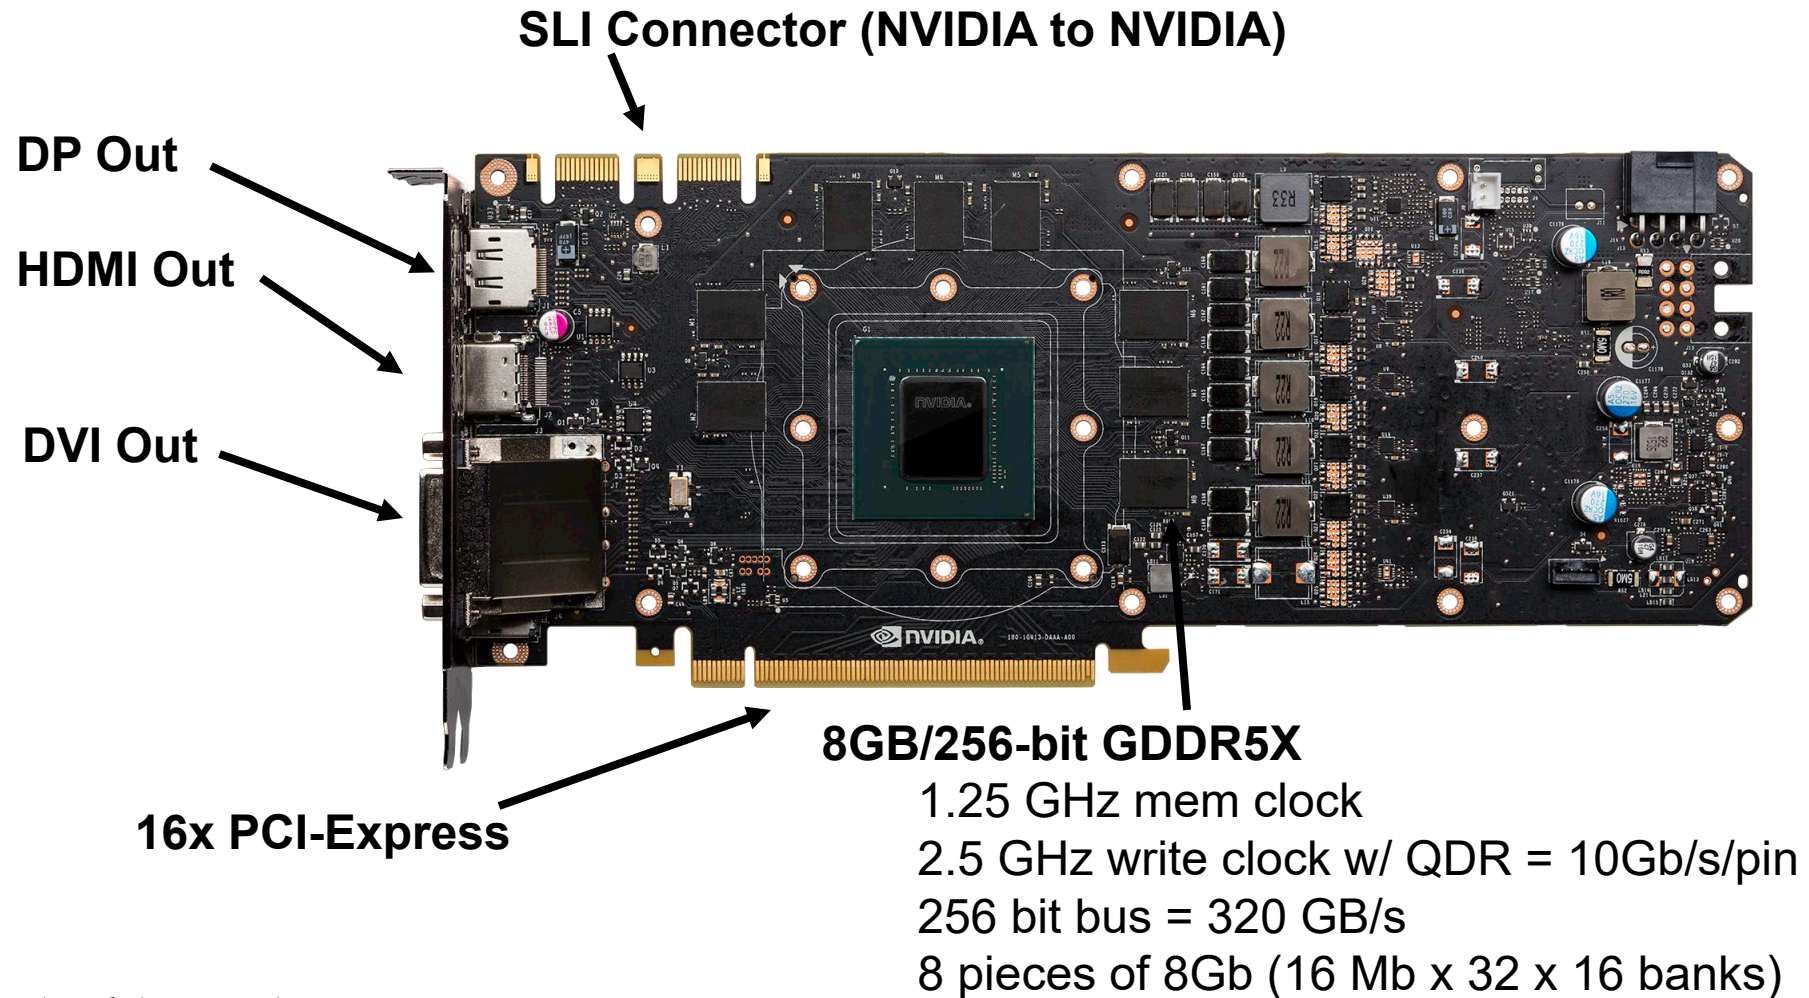
# Modern Intel PCIe PC Architecture



# Modern AMD PCIe PC Architecture



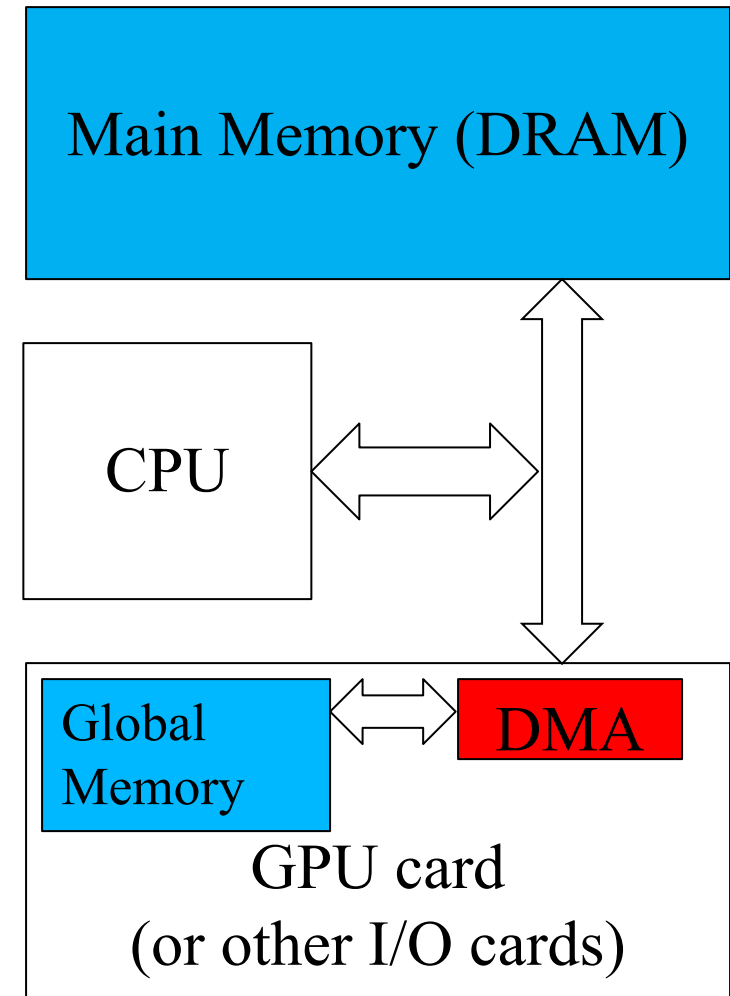
# GeForce GTX 1080 (Pascal) GPU Consumer Card Details



# PCIe Data Transfer using DMA

DMA (Direct Memory Access) is used to fully utilize the bandwidth of an I/O bus

- DMA uses physical address for source and destination
- Transfers a number of bytes requested by OS
- Needs pinned memory



# Pinned Memory

- DMA uses physical addresses
  - The OS could accidentally page out the data that is being read or written by a DMA and page in another virtual page into the same location
  - Pinned memory cannot be paged out
- If a source or destination of a `cudaMemcpy` in the host memory is not pinned, it needs to be first copied to a pinned memory – extra overhead
  - `cudaMemcpy` is much faster with pinned host memory source or destination

# Allocate/Free Pinned Memory (a.k.a. Page Locked Memory)

- `cudaHostAlloc()`
  - Three parameters
  - Address of pointer to the allocated memory
  - Size of the allocated memory in bytes
  - Option – use `cudaHostAllocDefault` for now
- `cudaFreeHost()`
  - One parameter
  - Pointer to the memory to be freed



# Using Pinned Memory

- Use the allocated memory and its pointer the same way those returned by `malloc()`;
- The only difference is that the allocated memory cannot be paged by the OS
- The `cudaMemcpy` function should be about 2x faster with pinned memory
- Pinned memory is a limited resource whose over-subscription can have serious consequences

# Important Trends

- Knowing yesterday, today, and tomorrow
  - The PC world is becoming flatter
  - CPU and GPU are being fused together
  - Outsourcing of computation is becoming easier...



Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

# ANY MORE QUESTIONS